Project Description Document

Numerical dataset

General Information on Dataset

- **Dataset Name:** **housing.csv**
- **Number of Classes:** NONE
- **Labels:** NONE
- **Total Number of Samples:** **20640**
- **Sample Size (if applicable):** [numerical]
- **Samples Used:**
    - Training: 16512
    - Validation: [none]
    - Testing: 4128

Implementation Details

- **Feature Extraction Phase:**
    - **Number of Features Extracted:** 13

Feature Names: **['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income', 'ocean_proximity_<1H OCEAN', 'ocean_proximity_INLAND', 'ocean_proximity_ISLAND', 'ocean_proximity_NEAR BAY', 'ocean_proximity_NEAR OCEAN']**

-

Dimension of Resulted Features: **(20640, 13)**

- **Cross-Validation:**
  - **Used?:** [**Yes**/No]
  - **Number of Folds:** [**5**]
  - **Training/Validation Ratio:** **4:1**
- **Hyperparameters Used:**
  - **Initial Learning Rate:** [Rate]
  - **Optimizer:** [Optimizer]
  - **Regularization:** [Type and strength]
  - **Batch Size:** [Size]
  - **Number of Epochs:** [Number]

Results Details

- ## Performance on Testing Data:

- **Linear Regression Score: 62.54**
- **Linear Regression MSE on Test Data: 4908476721.156613**
- **Linear Regression R-squared on Test Data: 0.6254240620553608**

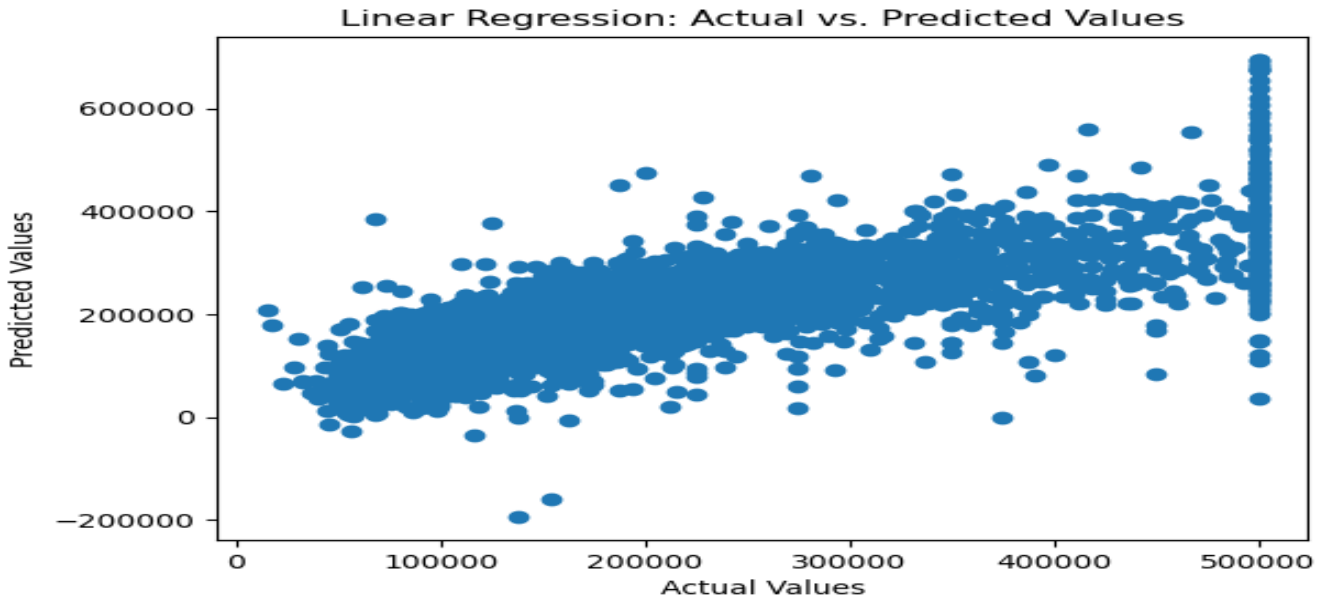

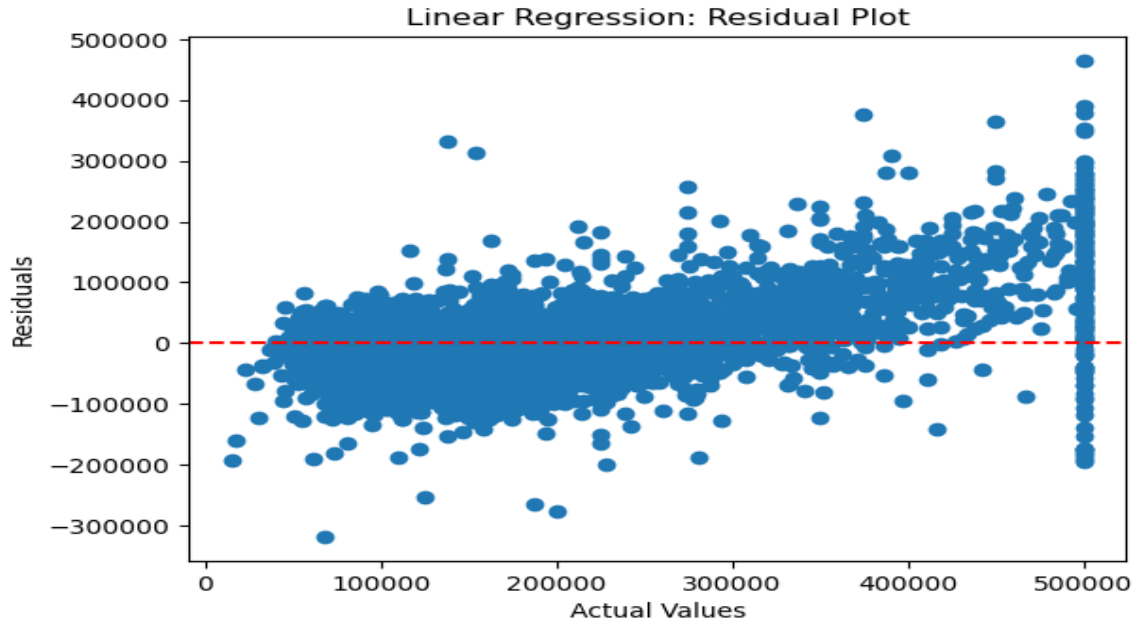

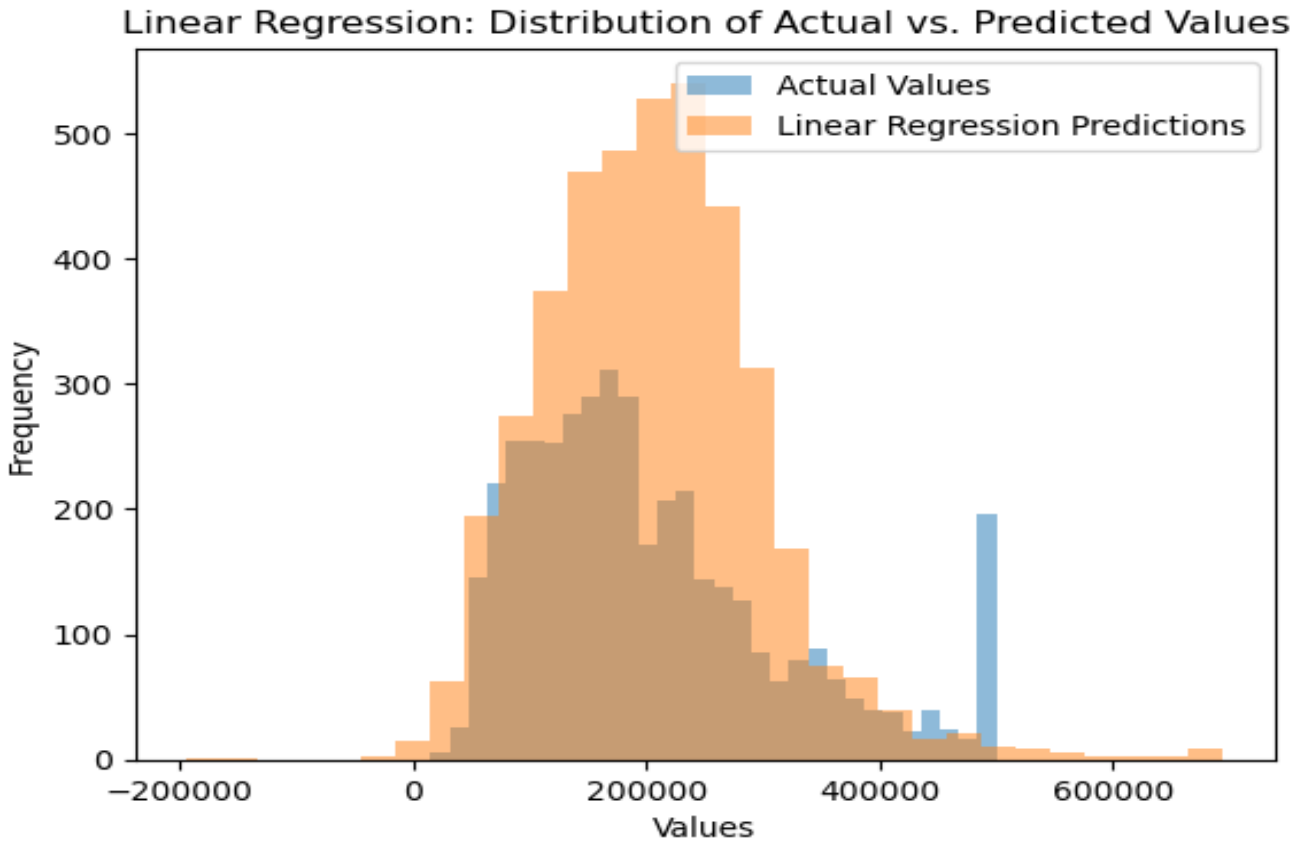
Linear Regression: Distribution of Actual vs. Predicted Values
Actual Values
Linear Regression Predictions
Frequency
Values
0
100
200
300
400
500
−200000
0
200000
400000
600000

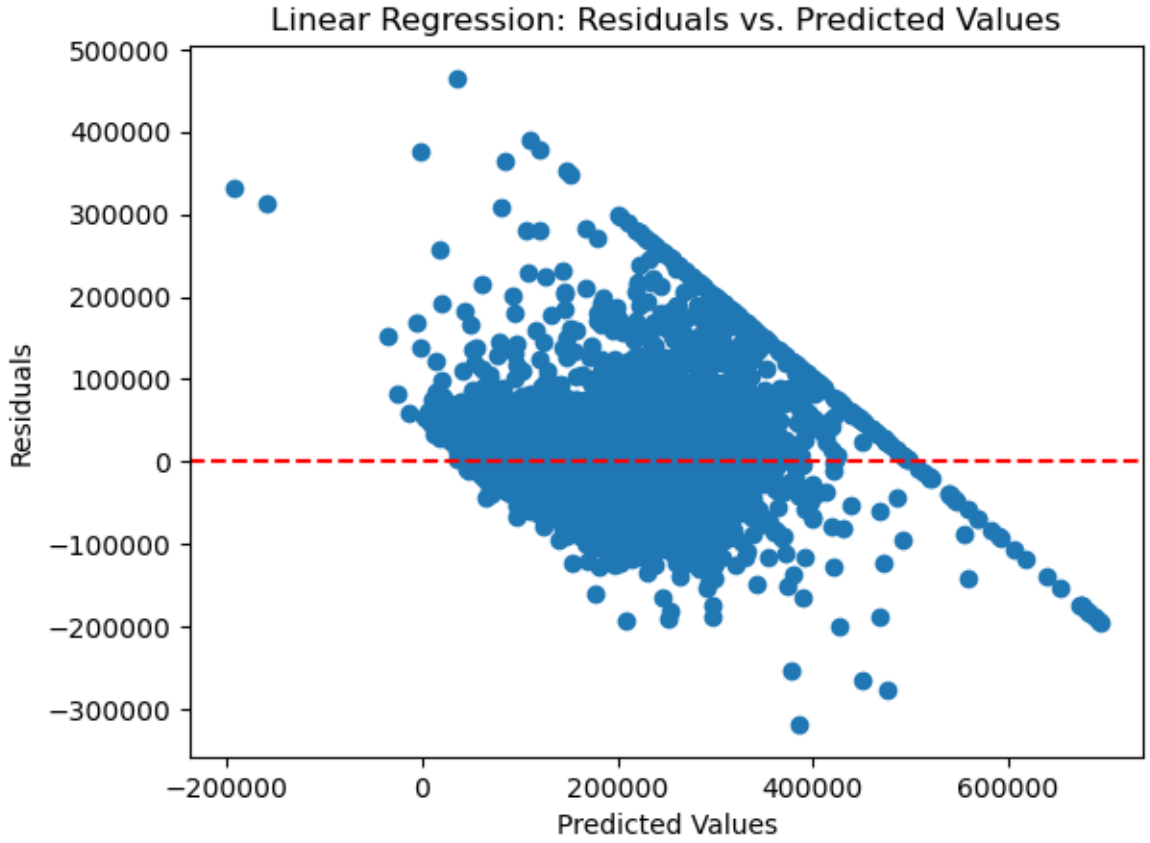
Linear Regression: Residuals vs. Predicted Values
Residuals
Predicted Values
500000
400000
300000
200000
100000
0
−100000
−200000
−300000
−200000
0
200000
400000
600000

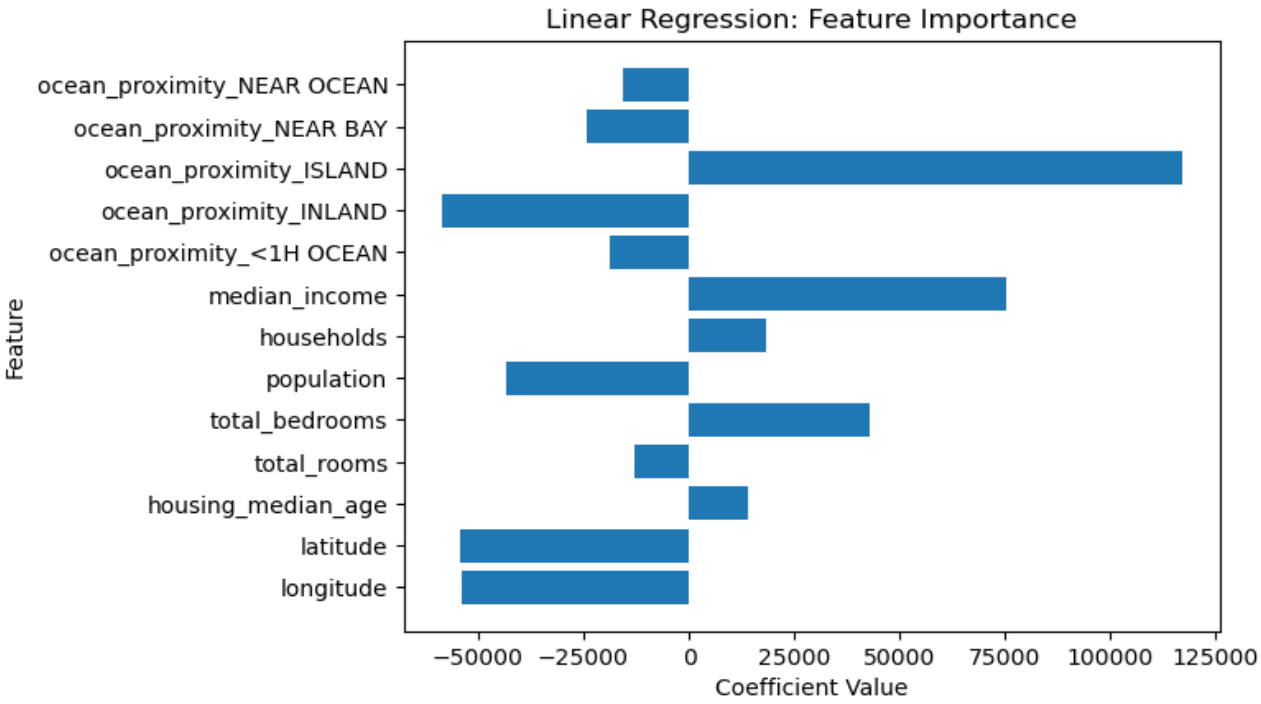
Linear Regression: Feature Importance
ocean_proximity_NEAR OCEAN
ocean_proximity_NEAR BAY
ocean_proximity_ISLAND
ocean_proximity_INLAND
ocean_proximity_<1H OCEAN
median_income
households
population
total_bedrooms
total_rooms
housing_median_age
latitude
longitude
Feature
−50000 −25000 0 25000 50000 75000 100000 125000
Coefficient Value

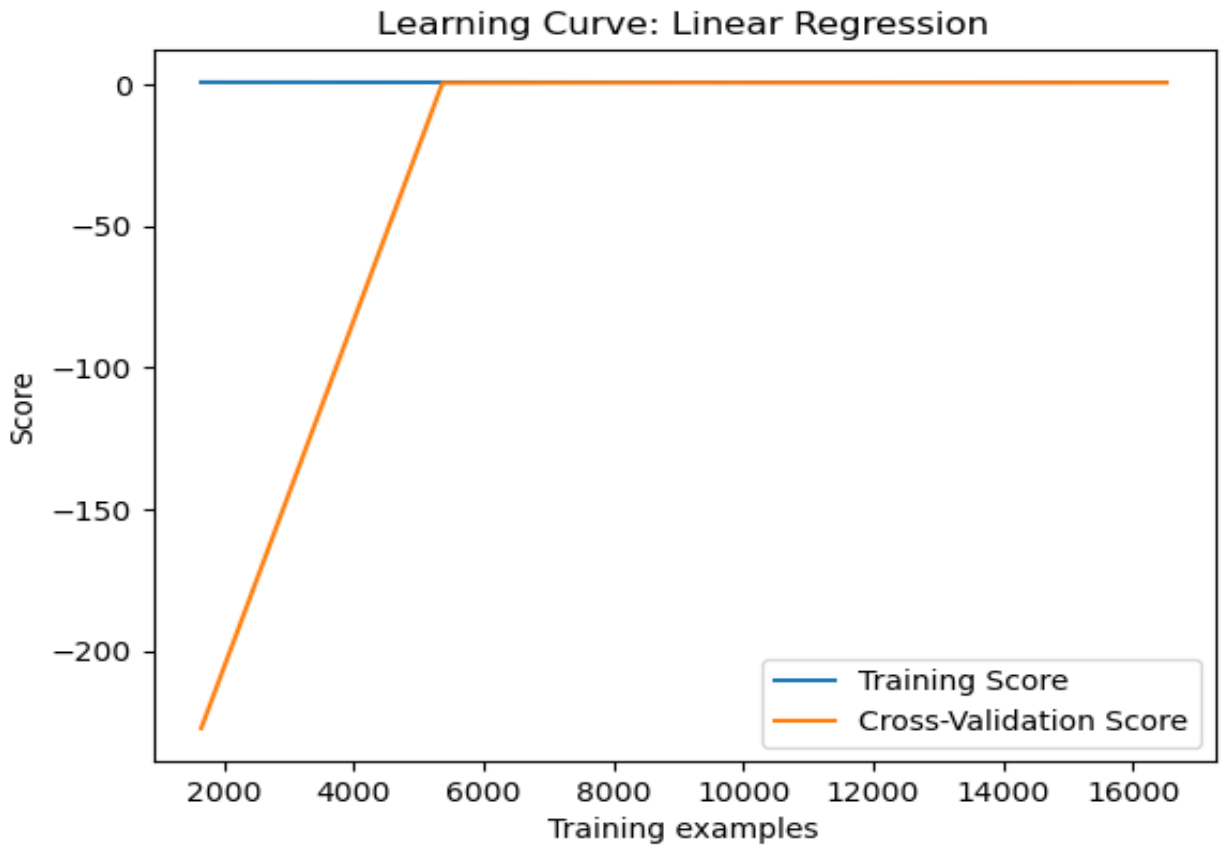
Learning Curve: Linear Regression
0
−50
−100
−150
−200
Score
2000 4000 6000 8000 10000 12000 14000 16000
Training examples
Training Score
Cross-Validation Score

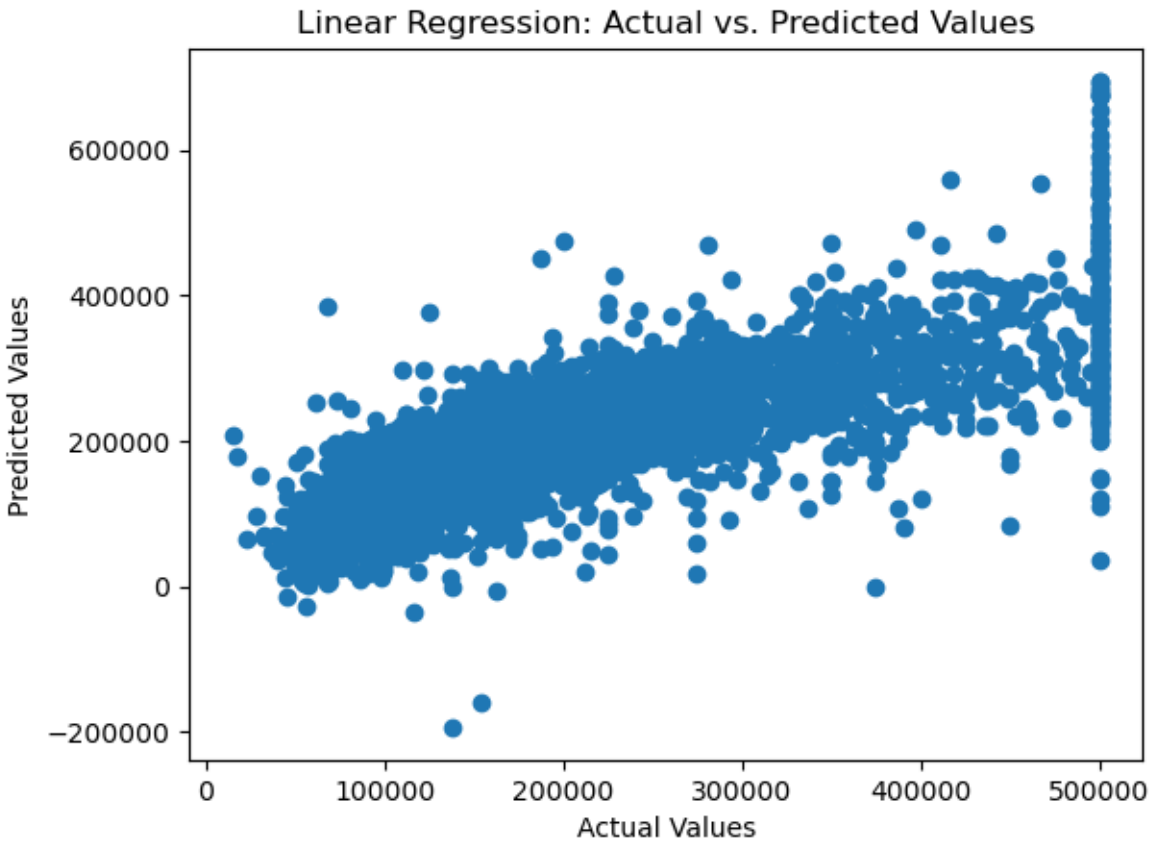
Linear Regression: Actual vs. Predicted Values
Predicted Values
600000
400000
200000
0
−200000
0
100000
200000
300000
400000
500000
Actual Values

KNN

- **Hyperparameters Used: K=5**

Results Details

- **Performance on Testing Data:**
- **KNN Regression Score: 71.3**
- **KNN Regression MSE on Test Data: 3760982284.460552**
- **KNN Regression R-squared on Test Data: 0.7129917188518262**

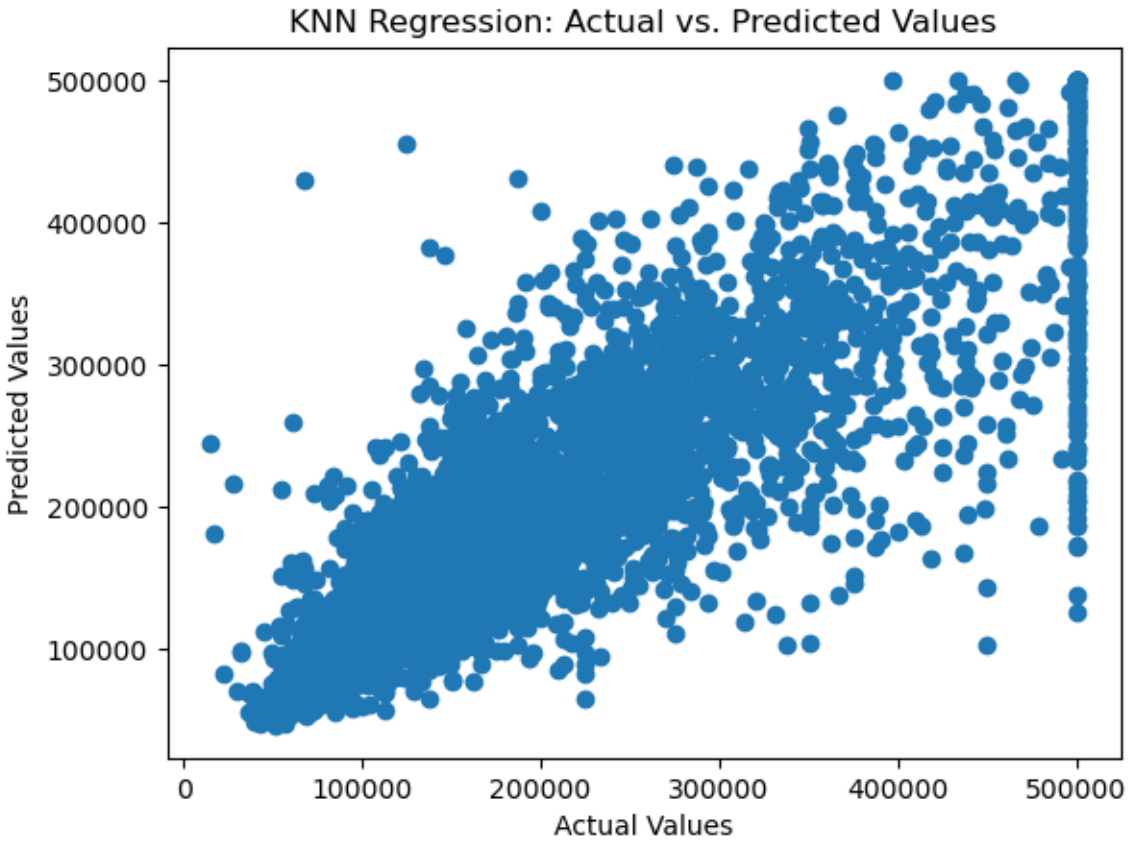

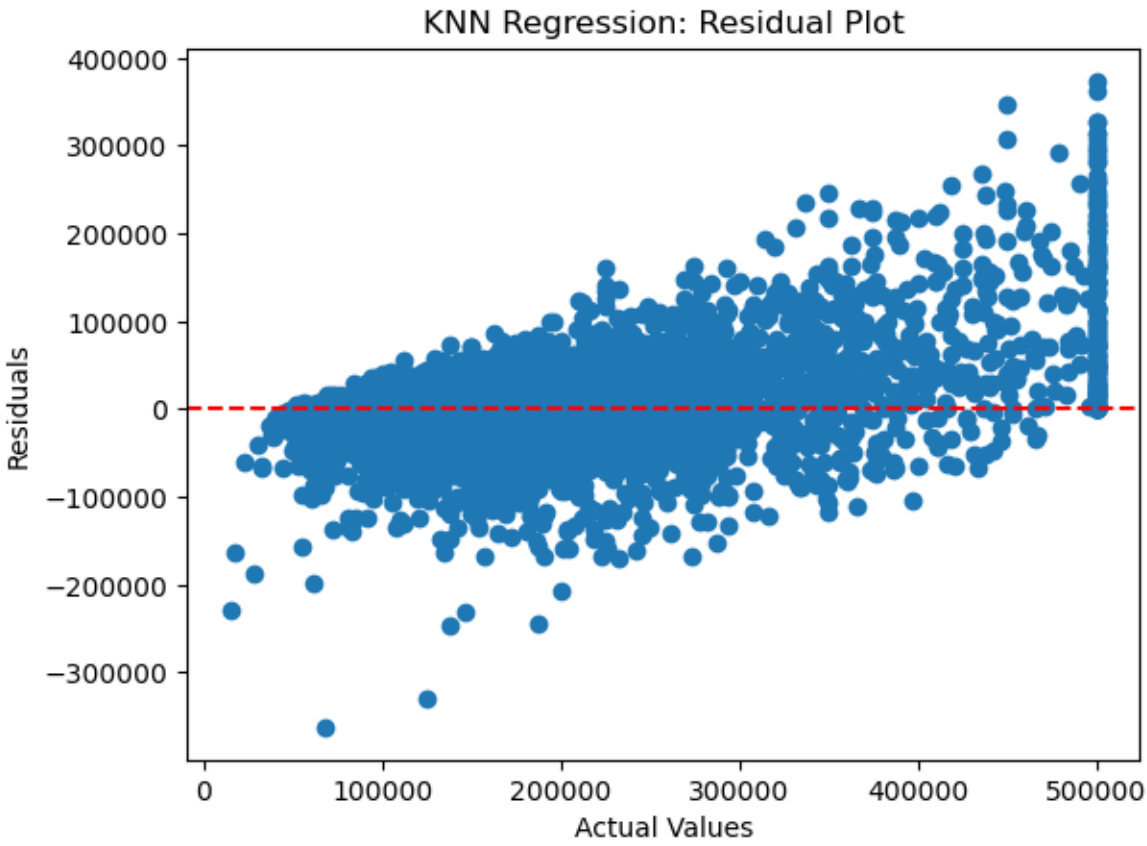
KNN Regression: Residual Plot
Residuals
Actual Values
400000
300000
200000
100000
0
−100000
−200000
−300000
0
100000
200000
300000
400000
500000

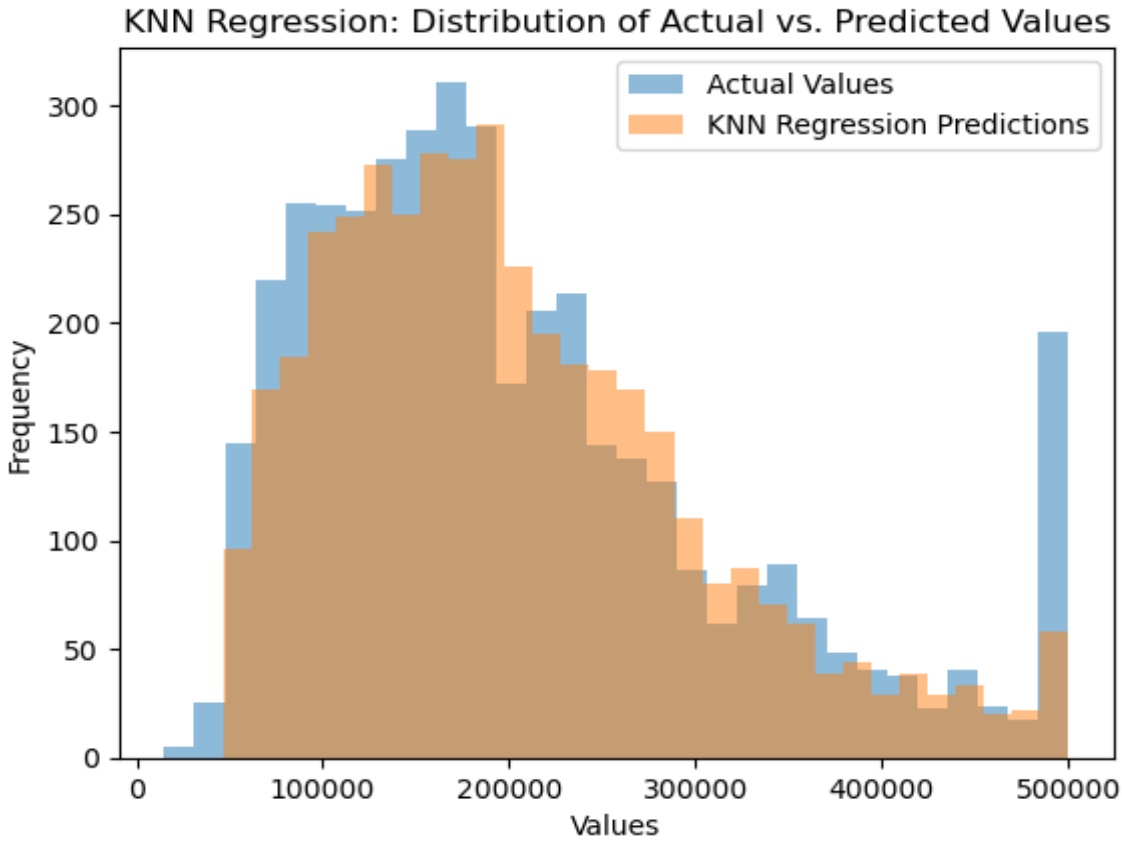
KNN Regression: Distribution of Actual vs. Predicted Values
Actual Values
KNN Regression Predictions
Frequency
Values
300
250
200
150
100
50
0
0
100000
200000
300000
400000
500000

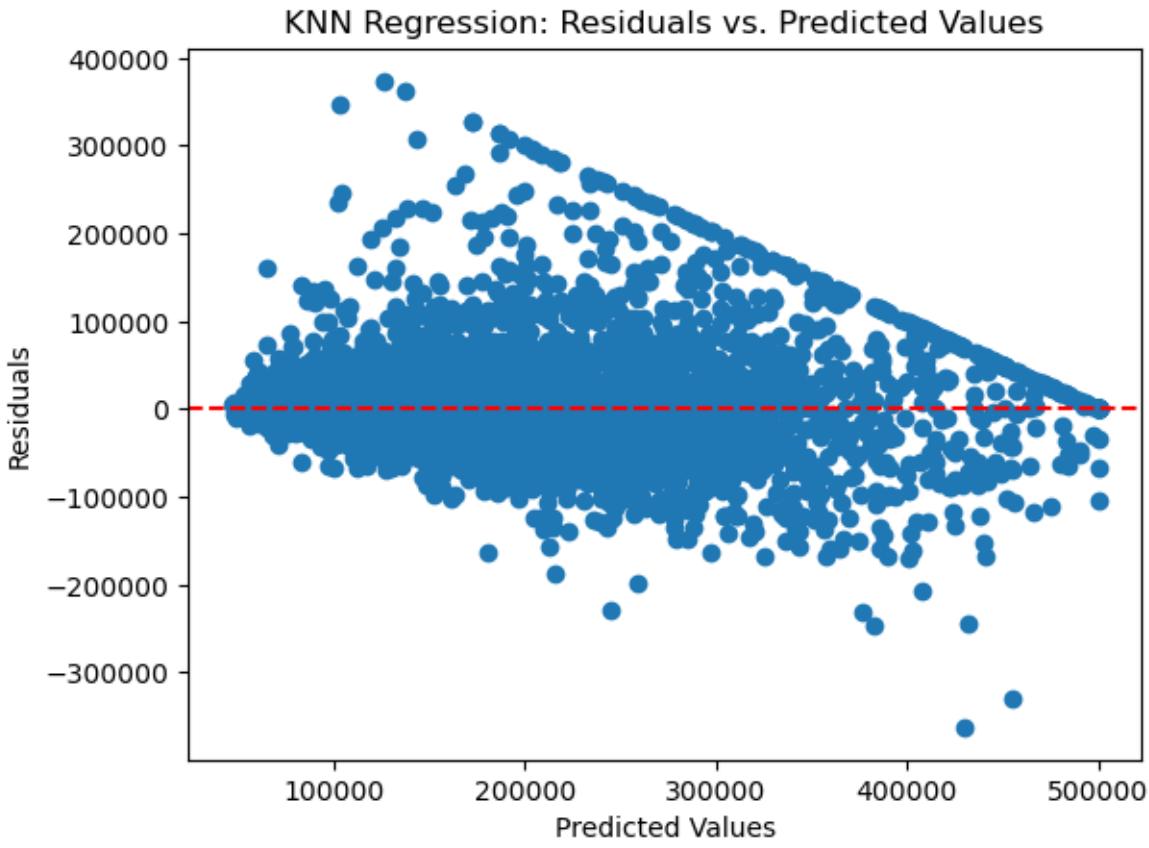
KNN Regression: Residuals vs. Predicted Values
Residuals
Predicted Values
400000
300000
200000
100000
0
−100000
−200000
−300000
100000
200000
300000
400000
500000

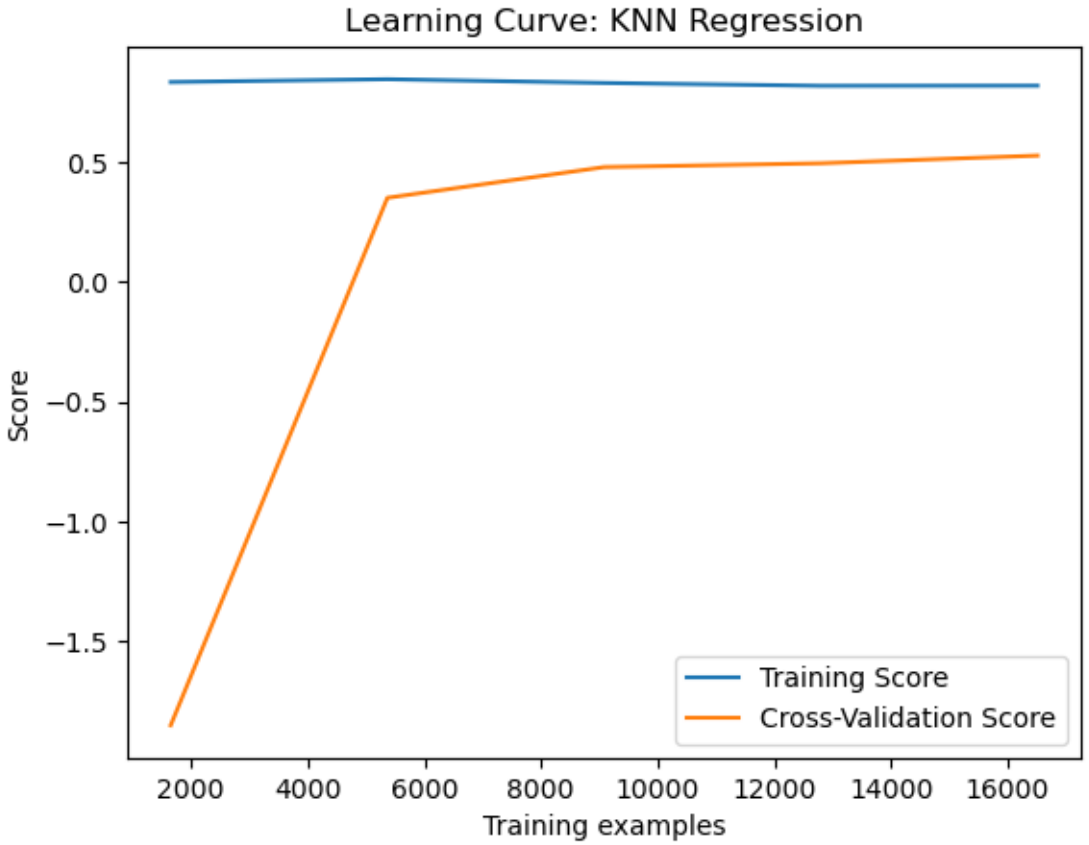
Learning Curve: KNN Regression
Score
Training examples
0.5
0.0
−0.5
−1.0
−1.5
2000
4000
6000
8000
10000
12000
14000
16000
Training Score
Cross-Validation Score

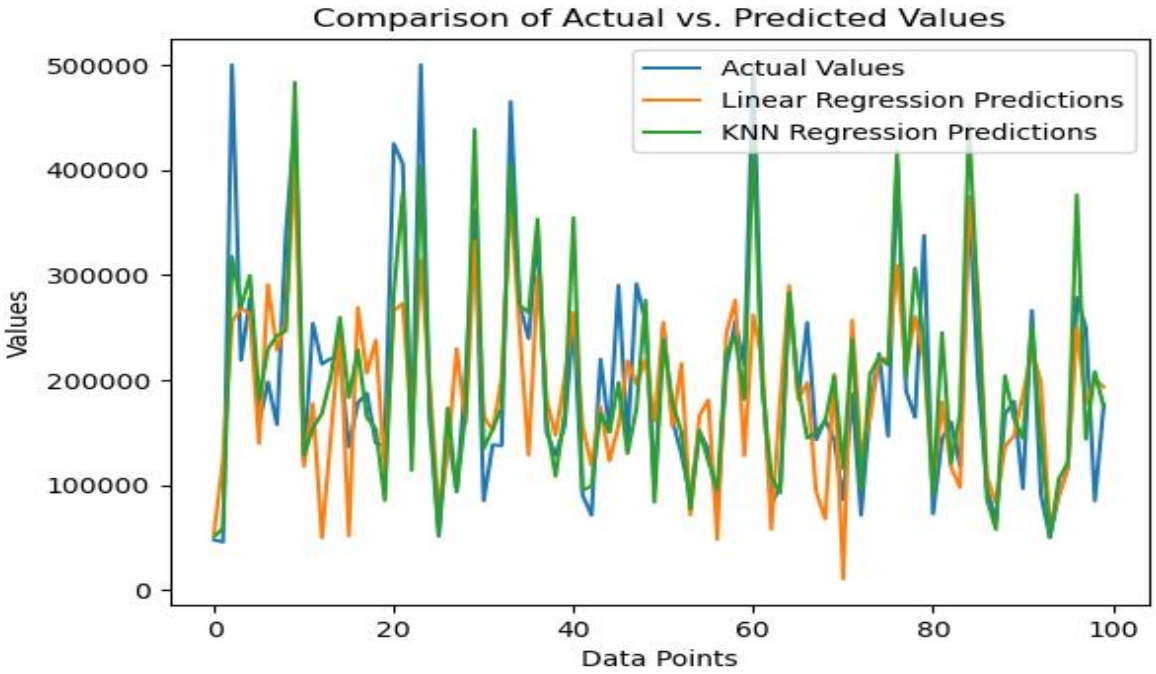
Comparison of Actual vs. Predicted Values
Actual Values
Linear Regression Predictions
KNN Regression Predictions
Values
Data Points
0
100000
200000
300000
400000
500000
0
20
40
60
80
100

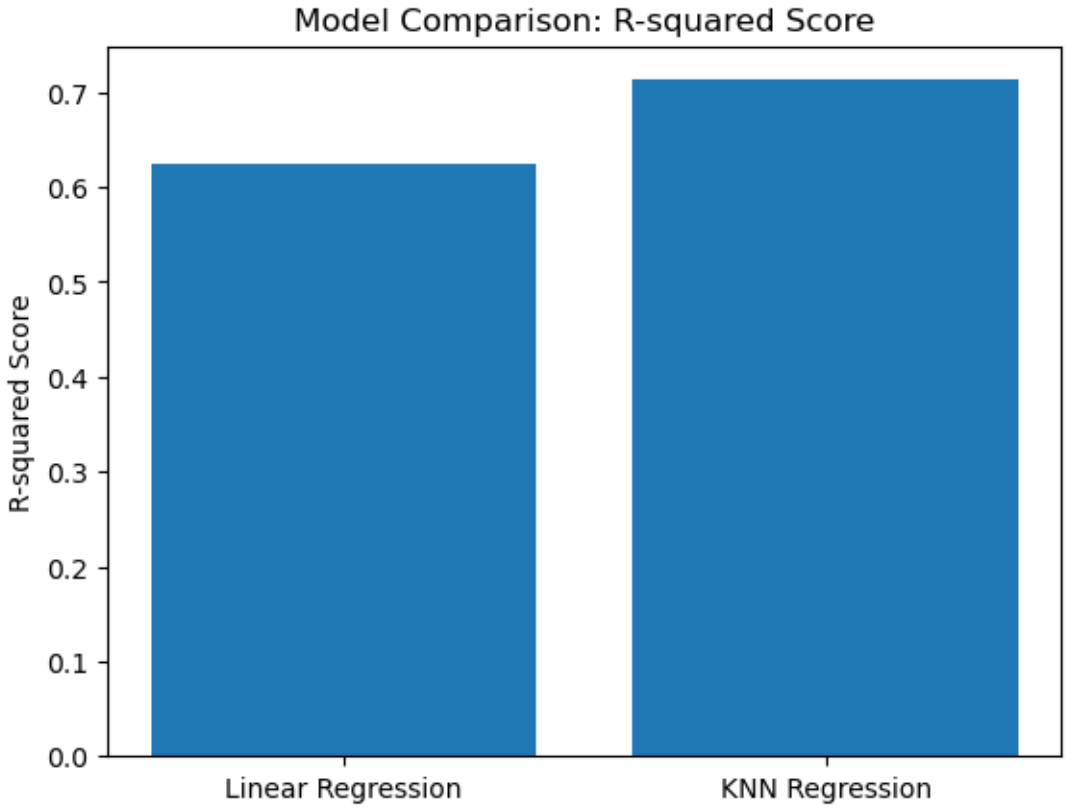
Model Comparison: R-squared Score
R-squared Score
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
Linear Regression
KNN Regression

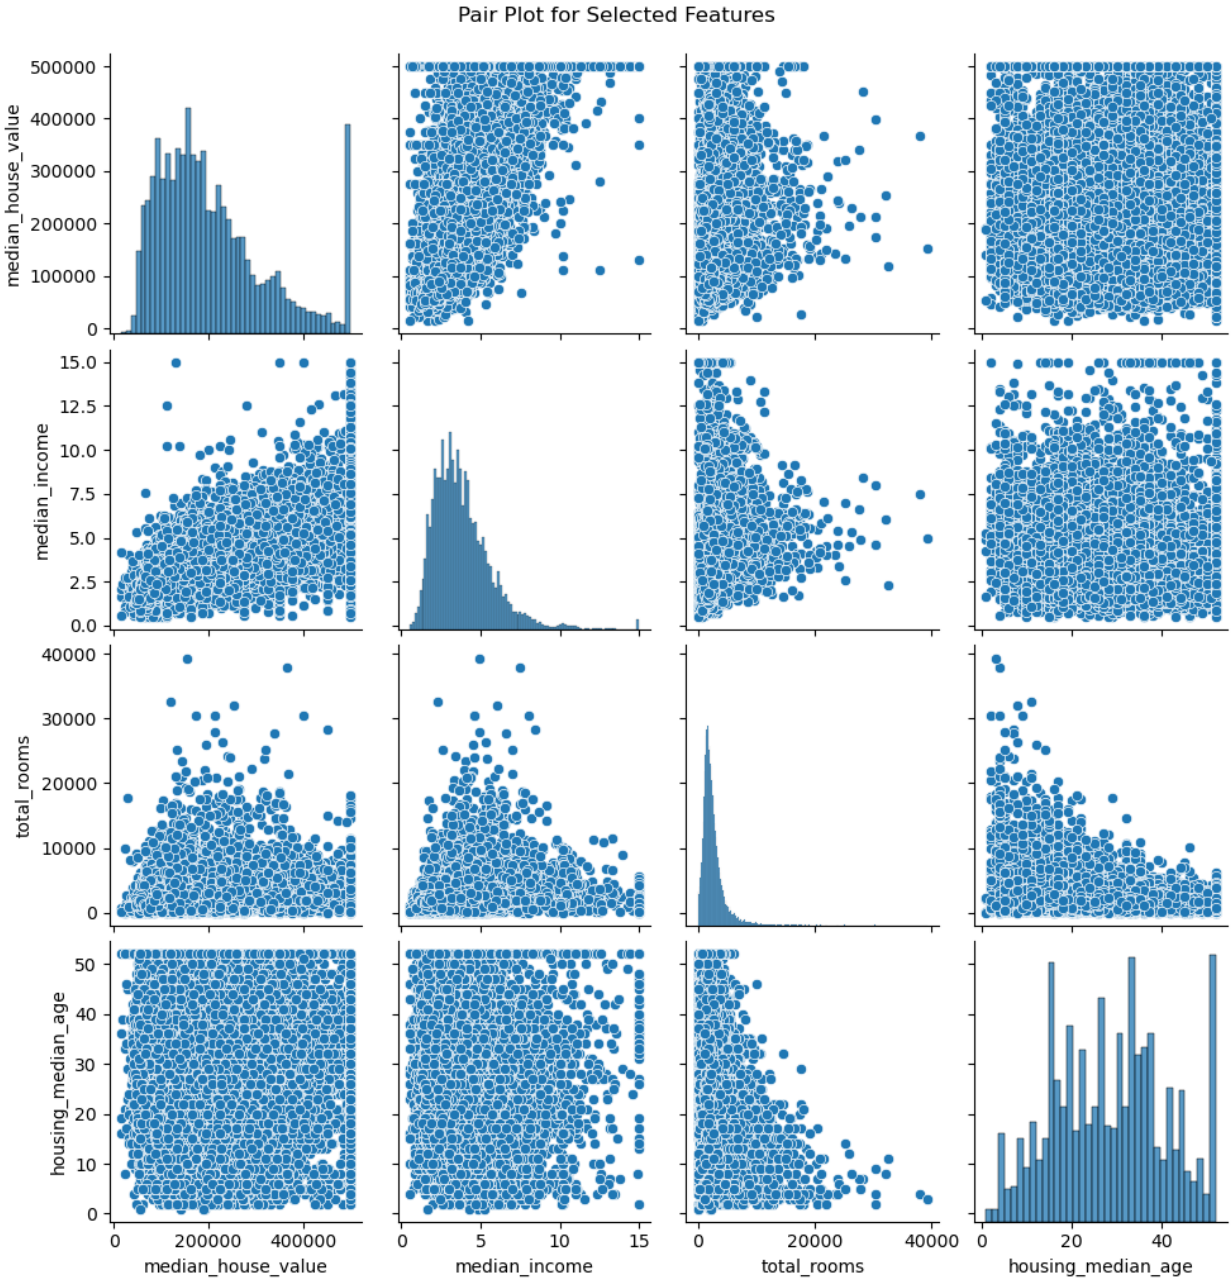
Pair Plot for Selected Features
median_house_value
median_income
total_rooms
housing_median_age
0
100000
200000
300000
400000
500000
0.0
2.5
5.0
7.5
10.0
12.5
15.0
0
10000
20000
30000
40000
0
10
20
30
40
50
0
200000
400000
0
5
10
15
0
20000
40000
0
20
40

Logistic Regression Model

Image dataset

General Information on Dataset

- **Dataset Name: dataset.csv**
- **Number of Classes: 3**
- **Labels: [0 2 3]**
- **Total Number of Samples: 35887**
- **Sample Size (if applicable): (48,48)**
- **Samples Used:**
  - **Training: 2400**
  - **Validation: 1800**
  - **Testing: 300**

Implementation Details

- **Feature Extraction Phase:**
- **Number of Features Extracted:**

Number of features extracted per image: 800

  - **Feature Names: HoG**
  - **Dimension of Resulted Features: (800,)**
- **Cross-Validation:**
  - **Used?:** [Yes/**No**]
  - **Number of Folds:** [Number]
  - **Training/Validation Ratio:** [Ratio]
  -

- **Hyperparameters Used:**
- **solver='lbfgs': The optimization algorithm to use. The 'lbfgs' solver is used for small datasets and is the default choice for Logistic Regression in scikit-learn.**
- **penalty='l2': The norm used in the penalization. 'l2' refers to the L2 regularization.**
- **C=1.0: Inverse of regularization strength; smaller values specify stronger regularization.**
- **max_iter=5000: Maximum number of iterations taken for the solver to converge.**

Results Details

- **Performance on Testing Data:**
- **Accuracy of Logistic Regression on HOG Features: 68%**

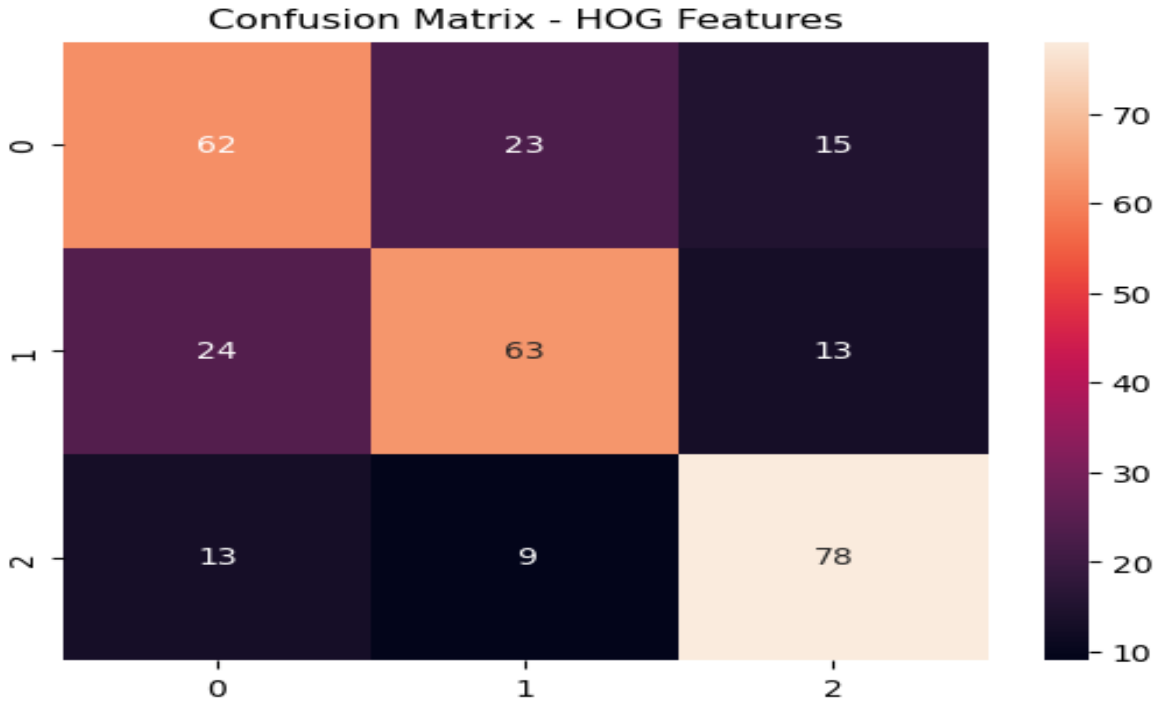

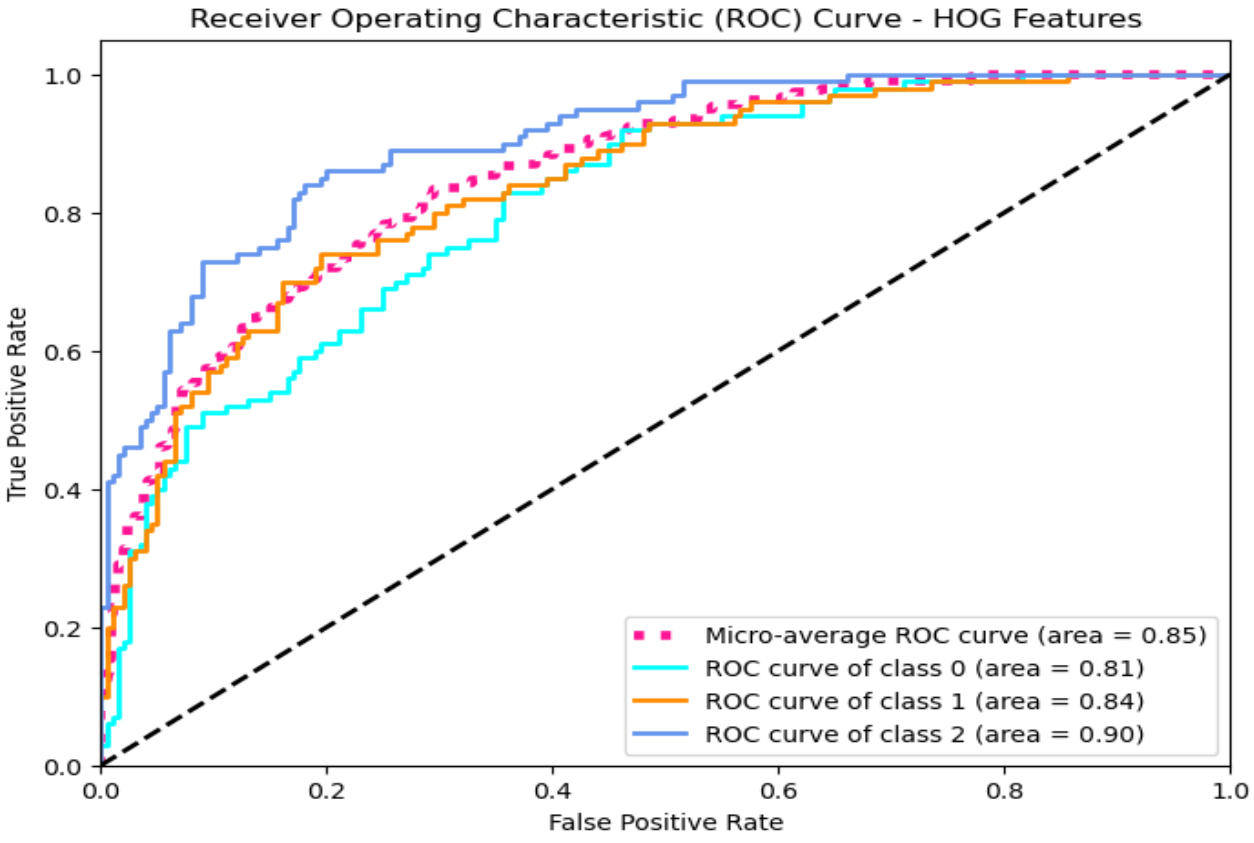

KMEANS MODEL

General Information on Dataset

- **Hyperparameters Used: k=3**

Results Details

- **Performance on Testing Data: 25%**

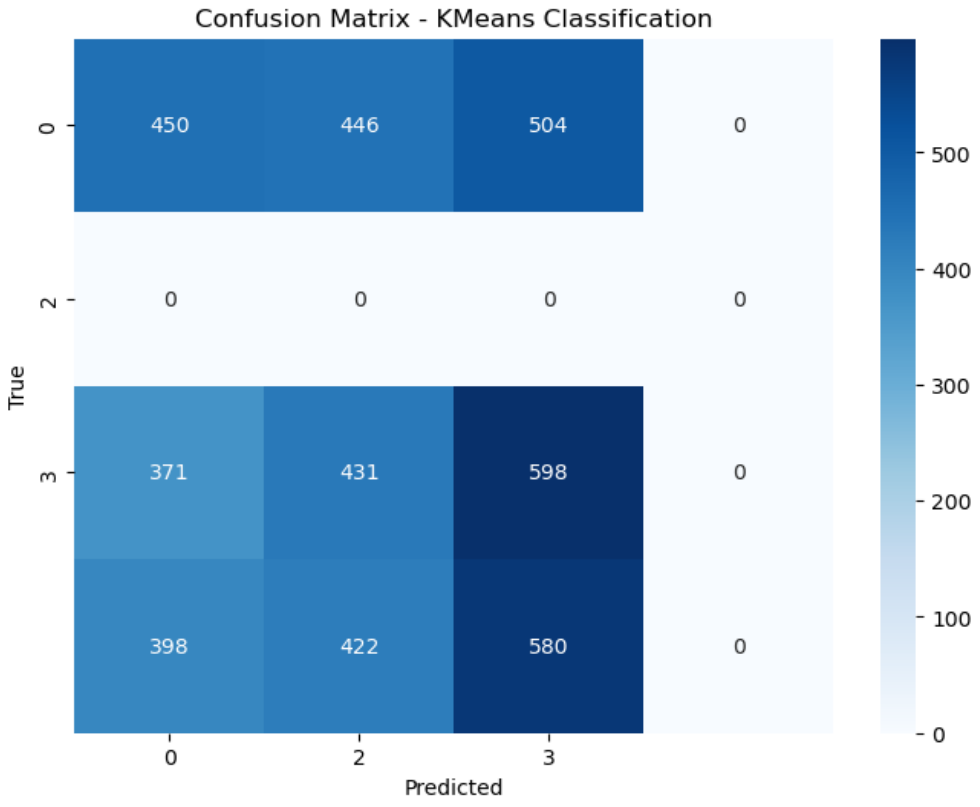

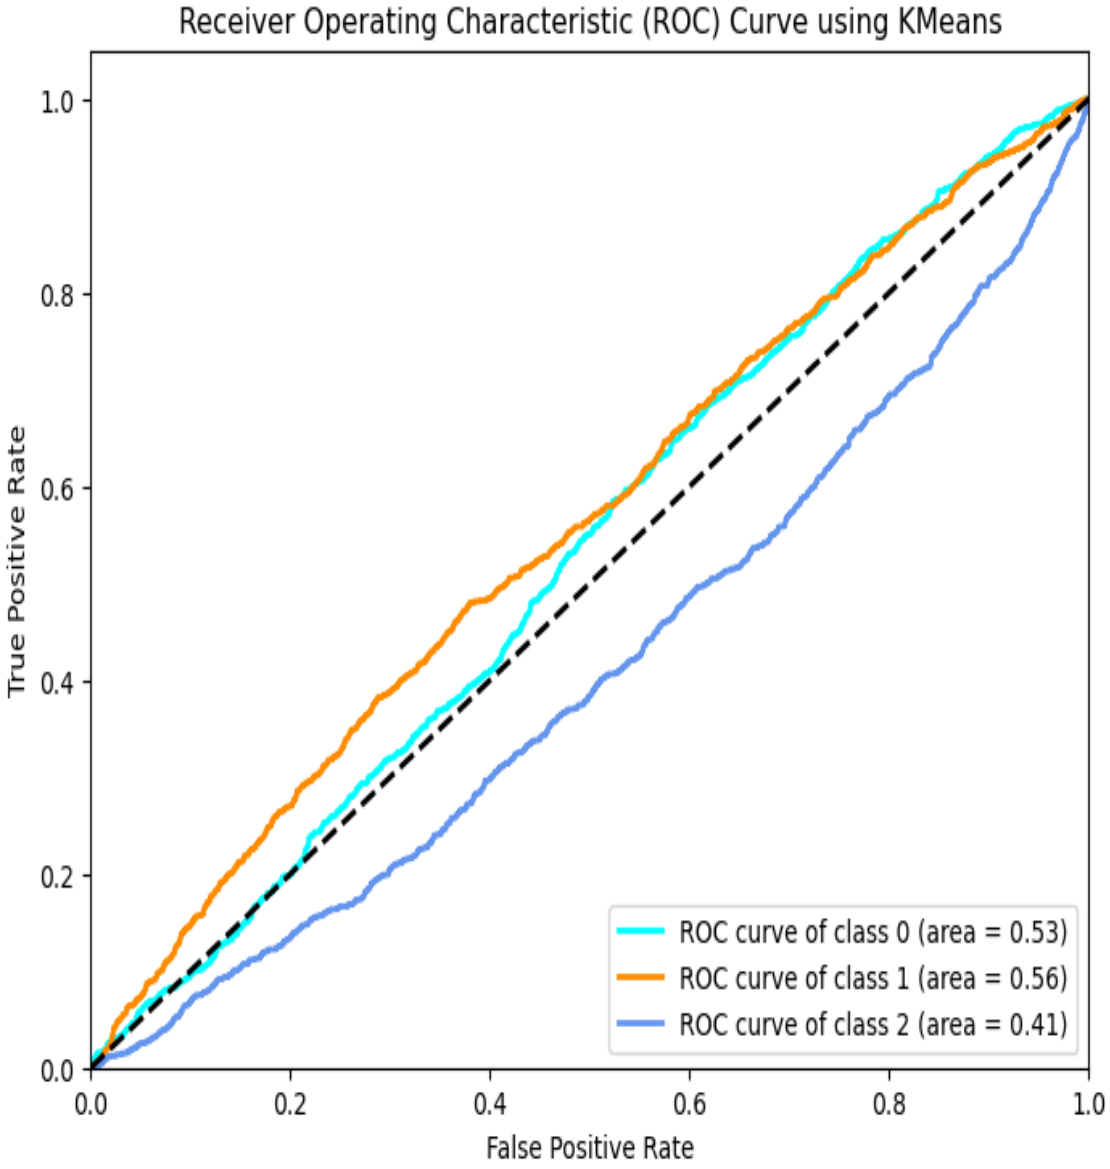
Receiver Operating Characteristic (ROC) Curve using KMeans
True Positive Rate
False Positive Rate
0.0
0.2
0.4
0.6
0.8
1.0
ROC curve of class 0 (area = 0.53)
ROC curve of class 1 (area = 0.56)
ROC curve of class 2 (area = 0.41)